



4 Courses

Fundamentals of Reinforcement Learning

Sample-based Learning Methods

Prediction and Control with Function Approximation

A Complete Reinforcement Learning System (Capstone)



06/23/2020

Yuki Kitayama

has successfully completed the online, non-credit Specialization

Reinforcement Learning

The Reinforcement Learning Specialization consists of 4 courses exploring the power of adaptive learning systems and artificial intelligence (AI). In this specialization, learners were taught to: Build a Reinforcement Learning system for sequential decision making; understand the space of Reinforcement Learning algorithms (Temporal- Difference learning, Monte Carlo, Sarsa, Q-learning, Policy Gradients, Dyna, and more); understand how to formalize a task as a Reinforcement Learning problem, and how to begin implementing a solution; understand how RL fits under the broader umbrella of machine learning. This learner is now prepared to take more advanced courses in AI or apply AI tools to real world problems.

Adam White
Assistant Professor
Computing Science
Faculty of Science

Martha White
Assistant Professor
Computing Science
Faculty of Science

The online specialization named in this certificate may draw on material from courses taught on-campus, but the included courses are not equivalent to on-campus courses. Participation in this online specialization does not constitute enrollment at this university. This certificate does not confer a University grade, course credit or degree, and it does not verify the identity of the learner.

Verify this certificate at:
coursera.org/verify/specialization/3HYEJUNANGZ2